

# Graph Neural Network–Driven EV Route Optimization from Live Traffic & Signal APIs

**Authors:** Aniket Gulab Khedkar

**Project Duration:** November 2024 – April 2024

**Date:** April 2025

## Abstract

We present an end-to-end pipeline that ingests **live traffic speeds** and **traffic-signal timing** from public/vendor APIs, cleans and fuses them into a **road graph**, and learns **edge-level travel time and stop-probability** with a **Graph Neural Network (GNN)**. The learned costs feed a **resource-constrained route optimizer** that respects **electric-vehicle (EV) battery** limits and can insert **charging stops** when necessary. The system builds a large-scale urban graph ( $\approx 114,883$  road/intersection nodes;  $\approx 500$  EV charging stations) and returns precise, real-time routes. Visual overlays mark EV stations (green), intersections (red), and roads (blue); routes show start (blue), end (red), and traversed path (white).

## 1. Introduction

Conventional shortest-path routing relies on static edge weights (length/speed-limit) or coarse historical averages. Urban

mobility, however, is **temporal, signal-controlled, and congested**.

Simultaneously, EV routing is **energy-constrained**, requiring **SoC-aware** path planning and **charging detours**. We address both problems by (i) learning **time-varying edge costs** (ETA and signal delay) with a GNN over the road graph, and (ii) solving a **constrained shortest path with resources** (battery, time) to deliver realistic, EV-feasible routes.

## 2. System Overview

**Data** → **Graph** → **GNN** → **Route Optimizer** → **API**

- Data ingestion (APIs):**
  - Link-level speeds/flow, incidents, and travel-time snapshots ( $\Delta t = 1\text{--}5$  min).
  - Signal programs (cycle length, phase splits, offsets, coordination groups, detector actuations).
  - Static base map (road geometry & attributes), elevation tiles (for grade).
- Preprocessing:** schema harmonization, deduplication, spatial snapping, time alignment, and **CSV exports** for reproducibility.
- Graph construction (NetworkX/GeoPandas):** directed multi-graph  $G=(V,E)$  with intersections & EV stations as nodes; road segments as edges with static and dynamic attributes.
- Learning: spatio-temporal GNN** to predict edge-time ETA and **signal stop probability**.
- Routing: A\*** with a fast heuristic on the **learned costs**, extended to a **resource-constrained label-**

setting algorithm for EV battery and optional charging stops.

---

## 3. Data & Pre-processing

### 3.1 Sources & schema

- **Traffic API:** link id, timestamp, speed, travel time, confidence.
- **Signals API:** intersection id, phase plan (C, g, y, r), offsets, coordination, current phase, queue length (if available).
- **Base map:** road centerlines, lanes, speed limit, oneway, functional class; node coords; segment length; curvature; **grade** derived from elevation rasters.
- **EV stations:** location, connector type, power (kW), access hours.

All feeds are normalized to UTC, deduplicated by (id, timestamp), and written to CSV staging tables.

### 3.2 Cleaning & fusion

- **Map-matching:** traffic link ids  $\rightarrow$  base-map edges; signals  $\rightarrow$  nearest node within  $\epsilon$  meters; resolve many-to-one cases via geometry overlap.
- **Temporal alignment:** resample all feeds to a common grid (e.g., 5-min buckets) with **last-observation-carried-forward** for  $\leq 10$ -min gaps; mark imputed flags.
- **Outlier handling:** Huber clipping on speeds; discard negative/zero travel times; enforce logical bounds (e.g.,  $v \leq 1.3 \times$  speed limit).
- **Feature standardization:** z-score for continuous features; one-hot for categorical (functional class, control type).

### 3.3 Graph statistics

The resulting city-scale graph contains  $\approx 114,883$  **intersection nodes** and  $\approx 500$  **EV charging stations** (marked as special nodes), with roads as directed edges; map renders use green (EV), red (intersections), blue (roads), and route overlays with white path from blue start to red end.

---

## 4. Graph Representation & Features

### 4.1 Nodes

- **Intersections:** degree, signal control type (fixed/actuated), offset within coordination band, historical stop rate.
- **EV stations:** charger power P<sub>kw</sub>, occupancy (if available), dwell-time distribution.

### 4.2 Edges

Each directed edge  $e=(u \rightarrow v)$  holds:

**Static:** length  $L_e$ , speed limit  $V_{\max}^{[f_0]}$ , lanes, grade  $\theta_e$ , curvature  $\kappa_e$ , functional class.

**Dynamic (per time slice  $t$ ):** speed  $v_{e,t}$ , travel time  $\tau_{e,t}$ , incident flags, **signal delay**  $\delta_{e,t}$  (expected red-time contribution), density proxy (e.g.,  $v/V_{\max}^{[f_0]}$ ), weather (optional).

**Target variables:**

- $\tau^{e,t}$  — predicted travel time (regression)
- $p^{e,t}_{\text{stop}}$  — probability of stopping at downstream signal (classification)

---

## 5. GNN Model

## 5.1 Architecture

We adopt a **spatio-temporal GraphSAGE** with a temporal encoder:

1. **Temporal encoder** (per edge): **TCN** (1D dilated conv) or **GRU** over the last  $W$  snapshots of edge features.
2. **Graph encoder** (per time  $t$ ): **GraphSAGE** / **GAT** on a **line-graph**  $L(G)$  (edges as nodes) to capture turn penalties & signal spillback, or message passing on  $G$  with **edge features** using PyTorch Geometric (NNConv, EdgeConv).
3. **Heads**:
  - Regression head  
→  $\tau^{e,t}$  (MSE loss)
  - Classification head  
→  $p^{e,tstop}$  (BCE or focal loss)

The joint loss:

$$L = \lambda \tau \text{MSE}(\tau_e, t, \tau^{e,t}) + \lambda p \text{BCE}(y_e, tstop, p^{e,tstop}) + \lambda \text{reg} \|\Theta\|_2^2.$$

## 5.2 Training protocol

- **Splits**: time-blocked (train: past weeks; val/test: future weeks) to avoid leakage; spatial holdout for robustness.
- **Sampling**: edge-time mini-batches with neighbor sampling  $k$  hops.
- **Optimization**: Adam (lr  $1e-3$ ), early stopping on **sMAPE** (for ETA) and **AUROC** (for stops).
- **Imbalance handling**: focal loss ( $\gamma=2$ ) and class weights for rare stop events.

## 6. EV Energy & Route Cost Modeling

### 6.1 Edge energy approximation

For edge  $e$  traversed at predicted speed  $v$  and grade  $\theta$ :

$$E_e \approx [c_0 + c_1 v + c_2 v^2 + c_3 \sin[f_0](\theta)] \cdot L_e,$$

where coefficients summarize **aero drag** ( $\propto v^2$ ), **rolling resistance**, and **grade**; regenerative braking is modeled via efficiency  $\eta_{\text{regen}}$  when  $\sin[f_0](\theta) < 0$ .

### 6.2 Multi-objective route cost

We minimize a convex combination:

$$C(\pi) = \sum_{e \in \pi} (w_\tau \tau^{e,t} + w_\delta p^{e,tstop} \delta^- + w_E E_e + w_{\text{turn}} \text{Penalty}_e),$$

subject to **battery constraints**:

$$\text{SoC}_{k+1} = \text{SoC}_k - E_e B_{\text{cap}} \text{ and } \text{SoC}_k \geq \text{SoC}_{\text{min}}.$$

Charging at station node  $s$  adds dwell time  $T_{\text{chg}}(s, \Delta \text{SoC})$  and increases  $\text{SoC}$  accordingly.

## 7. Routing Algorithms

### 7.1 Learned-cost $A^*$

We run  $A^*$  using a consistent heuristic  $h(n) = \text{euclidean}(n, \text{goal}) / V_{\text{free}}$ . Edge weights are the **GNN-predicted** travel time plus expected signal penalty and (optionally) an energy term weighted into equivalent seconds.

### 7.2 Resource-constrained shortest path (battery)

We solve a **Shortest Path Problem with Resource Constraints (SPPRC)** using a **label-setting** algorithm over states (node,  $\text{SoC}$ ):

```

Initialize label at (source, SoC0)
with cost 0
while queue not empty:
    pop best label (n, soc)
    for each edge e = (n -> m):
        cost' = cost + C_e
        soc' = soc - E_e/Bcap
        if soc' < SoCmin: continue
        if m is EV station and soc' <
SoC_thresh:
            consider charging options:
            soc'' = min(1.0, soc' + Δsoc),
            cost'' = cost' + Tchg
            relax (m, soc'') if
Pareto-better
            relax (m, soc') if Pareto-
better
Return best label at (goal, soc >=
SoCmin)

```

Dominance rules prune inferior labels on (cost, SoC); optional **bi-directional** expansion accelerates search.

## 8. Evaluation

### 8.1 Offline predictive metrics

- **ETA:** MAE / RMSE / sMAPE over edges & paths; calibration plots of predicted vs. observed times.
- **Stops:** AUROC / AUPRC and **Expected Red Time** accuracy at signalized approaches.

### 8.2 Routing quality metrics

- **Optimality gap:** (Cours–Coracle)/Coracle vs. a hindsight oracle (using realized times).
- **On-time arrival rate** under arrival deadlines.
- **Energy feasibility:** % routes meeting SoC constraints without emergency detours.
- **Charging efficiency:** added minutes vs. fixed-stop heuristics; station congestion sensitivity.

### Baselines:

- Static Dijkstra (free-flow speeds)
- Historical-mean weights
- Commercial-style ETA without signal model (if available)

### 8.3 Ablations

- **No-traffic / no-signal** ablations to quantify each source's lift.
- **GNN variants:** GraphSAGE vs. GAT vs. line-graph message passing.
- **Temporal window W** and resampling interval sensitivity.
- **Cost weights**  $w_\tau, w_\delta, w_E$  trade-off curves (Pareto front).

## 9. Implementation Details

- **Graph & geo:** Python, NetworkX, GeoPandas, Shapely; spatial index (rtree); optional PostGIS for persistence.
- **Learning:** PyTorch Geometric / DGL; mixed precision; neighbor sampling; experiment tracking with MLflow/W&B.
- **Pipelines:** pydantic configs; Airflow or Prefect for scheduled data fetch & model refresh.
- **Serving:** FastAPI endpoint `POST /route` accepting {source, dest, SoC0, SoCmin, timestamp}; GNN scoring cached per time slice; A\* in Cython/NumPy for speed.
- **Visualization:** Folium/Kepler.gl map layers (EV stations, intersections, edges, path).

## 10. Limitations & Risks

- **Sensor/API drift:** provider calibrations change; monitor residuals and retrain triggers.
- **Actuated signals:** stochastic phase times; we approximate with expected red; queue spillback may violate independence assumptions.
- **Energy model:** simplified coefficients; vehicle-specific parameters (mass, CdA, tires) should be profiled for accuracy.
- **Data bias:** coverage gaps in speed feeds; EV-station availability may be stale—expose uncertainty.

---

## 11. Reproducibility

- Version-pin dependencies; save **graph snapshots** (Parquet/Feather) per time slice.
- Log seeds, model state-dict, and train/val/test indices.
- Provide  
scripts: `fetch_apis.py` → `fuse_t  
o_csv.py` → `build_graph.py` →  
`train_gnn.py` → `serve_route.p  
y`.
- Ship a **synthetic toy city** for unit tests of SPPRC and charging insertion logic.

---

## 12. Conclusion

By fusing **traffic** and **signal timing** into a **graph learning** framework, then routing with **resource constraints**, we obtain routes that are **faster**, **more predictable**, and **EV-feasible**. The approach is modular—additional layers (incidents, weather, curb regulations) can be appended—and deployable in real time through cached GNN scores and accelerated search.

---

## Appendix A — Key Hyperparameters (example)

- Temporal window  $W=6$  (past 30 min at 5-min resolution), GraphSAGE depth=2, hidden=128, dropout=0.2
- Loss weights:  $\lambda\tau=1.0, \lambda p=0.5$ ; cost weights:  $w\tau=1, w\delta=0.5, wE=0.2, w_r=0.05$
- A\* heuristic speed  $V_{free}=60$  km/h; SoCmin = 10%; charge threshold 20%

## Appendix B — Data Schemas (simplified)

**traffic.csv:** `link_id, timestamp, speed, travel_time, conf`

**signals.csv:** `int_id, timestamp, cycle, offset, phase, split, control_type`

**stations.csv:** `node_id, lat, lon, kw, connectors`

**graph\_edges.csv:** `u, v, length, grade, lanes, func_class, speed_limit, oneway`

---